

# **Data Cleaning and Preparation**

## **Book**

*Python for Data Analysis*

O'REILLY®

2nd Edition

# Python for Data Analysis

DATA WRANGLING WITH PANDAS,  
NUMPY, AND IPYTHON



powered by



Wes McKinney

[www.allitebooks.com](http://www.allitebooks.com)

## CH 7 : Data Cleaning and Preparation

### chapter introduction

During the course of doing data analysis and modeling, a significant amount of time is spent on data preparation: loading, cleaning, transforming, and rearranging. Such tasks are often reported to take up 80% or more of an analyst's time. Sometimes the way that data is

stored in files or databases is not in the right format for a particular task. Many researchers choose to do ad hoc processing of data from one form to another using a general-purpose programming language, like Python, Perl, R, or Java, or Unix text-processing tools like sed or awk. Fortunately, pandas, along with the built-in Python language features, provides you with a high-level, flexible, and fast set of tools to enable you to manipulate data into the right form.

If you identify a type of data manipulation that isn't anywhere in this book or elsewhere in the pandas library, feel free to share your use case on one of the Python mailing lists or on the pandas GitHub site. Indeed, much of the design and implementation of pandas has been driven by the needs of real-world applications.

In this chapter I discuss tools for missing data, duplicate data, string manipulation, and some other analytical data transformations. In the next chapter, I focus on combining and rearranging datasets in various ways.

```
In [6]: import numpy as np
        from numpy import nan as NA
        import pandas as pd
```

## Handling Missing Data

```
In [2]: dt = pd.Series([np.nan,1,25,np.nan,60,99])
```

```
In [3]: dt
```

```
Out[3]: 0      NaN
        1      1.0
        2     25.0
        3      NaN
        4     60.0
        5     99.0
        dtype: float64
```

```
In [4]: dt.dropna()
```

```
Out[4]: 1      1.0
        2     25.0
        4     60.0
        5     99.0
        dtype: float64
```

```
In [5]: dt[dt.notnull()]
```

```
Out[5]: 1      1.0
        2     25.0
        4     60.0
        5     99.0
        dtype: float64
```

```
In [6]: data = pd.DataFrame([[1., 6.5, 3.], [1., NA, NA], [NA, NA, NA], [NA, 6.5, 3.]])
```

```
In [7]: data
```

```
Out[7]:
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
2	NaN	NaN	NaN
3	NaN	6.5	3.0

**if we passing how = 'all' to drop func, it 'll remove all rows that are all NA**

```
In [8]: data.dropna(how = 'all')
```

```
Out[8]:
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
3	NaN	6.5	3.0

**if we passing how = 'any' to drop func, it 'll remove all rows that have at least NA**

```
In [9]: data.dropna(how = 'any')
```

```
Out[9]:
```

	0	1	2
0	1.0	6.5	3.0

```
In [10]: data[4] = NA
         data[5] = 5.0
```

```
In [11]: data
```

```
Out[11]:
```

	0	1	2	4	5
0	1.0	6.5	3.0	NaN	5.0
1	1.0	NaN	NaN	NaN	5.0
2	NaN	NaN	NaN	NaN	5.0
3	NaN	6.5	3.0	NaN	5.0

**we can do the same way to drop columns by passing axis = 1**

```
In [12]: data.dropna(axis = 1, how = 'all')
```

```
Out[12]:
```

	0	1	2	5
0	1.0	6.5	3.0	5.0
1	1.0	NaN	NaN	5.0
2	NaN	NaN	NaN	5.0
3	NaN	6.5	3.0	5.0

```
In [13]: data.dropna(axis = 1, how = 'any')
```

```
Out[13]:
```

	5
0	5.0
1	5.0
2	5.0
3	5.0

```
In [ ]: data = data.drop(6, axis = 1)
```

**If you want to keep only rows containing a certain number of observations. You can indicate this with the thresh argument**

```
In [ ]: data.dropna(thresh=2)
```

```
In [ ]: data
```

```
In [ ]: [data.dropna(thresh=i) for i in range(1,6)]
```

## Filling In Missing Data

to fill NA values, we can replace them with real values:

```
In [ ]: data.fillna(0)
```

you can use a different fill value for each column:

```
In [16]: data.fillna({0:0, 1: 1, 2: 2, 3:3, 4:4 , 5:5})
```

```
Out[16]:
```

	0	1	2	4	5
0	1.0	6.5	3.0	4.0	5.0
1	1.0	1.0	2.0	4.0	5.0
2	0.0	1.0	2.0	4.0	5.0
3	0.0	6.5	3.0	4.0	5.0

```
In [15]: data.fillna(data.mean())
```

```
Out[15]:
```

	0	1	2	4	5
0	1.0	6.5	3.0	NaN	5.0
1	1.0	6.5	3.0	NaN	5.0
2	1.0	6.5	3.0	NaN	5.0
3	1.0	6.5	3.0	NaN	5.0

## Removing Duplicates

```
In [17]: data = pd.DataFrame({'k1': ['one', 'two'] * 3 + ['two'], 'k2': [1, 1, 2, 3, 3, 4,
```

```
In [18]: data
```

```
Out[18]:
```

	k1	k2
0	one	1
1	two	1
2	one	2
3	two	3
4	one	3
5	two	4
6	two	4

```
In [19]: data.duplicated()
```

```
Out[19]: 0    False
          1    False
          2    False
          3    False
          4    False
          5    False
          6     True
          dtype: bool
```

**drop\_duplicates** returns a DataFrame where the duplicated array is False

```
In [20]: data.drop_duplicates()
```

```
Out[20]:
```

	k1	k2
0	one	1
1	two	1
2	one	2
3	two	3
4	one	3
5	two	4

**Both of these methods by default consider all of the columns; alternatively, you can specify any subset of them to detect duplicates**

```
In [23]: data['v1'] = range(7)
```

```
In [24]: data
```

```
Out[24]:
```

	k1	k2	v1
0	one	1	0
1	two	1	1
2	one	2	2
3	two	3	3
4	one	3	4
5	two	4	5
6	two	4	6

```
In [25]: data.drop_duplicates(['k1'])
```

```
Out[25]:
```

	k1	k2	v1
0	one	1	0
1	two	1	1

**drop\_duplicates** and **drop\_duplicates** by default keep the first observed value combination. Passing **keep='last'** will return the last one:

```
In [26]: data.drop_duplicates(['k1', 'k2'], keep='last')
```

```
Out[26]:
```

	k1	k2	v1
0	one	1	0
1	two	1	1
2	one	2	2
3	two	3	3
4	one	3	4
6	two	4	6

## Mapping

- For many datasets, you may wish to perform some transformation based on the values in an array, Series, or column in a DataFrame.
- we use map function to perform like this transformation
- Using map is a convenient way to perform element-wise transformations and other data cleaning-related operations.

```
In [27]: data = pd.DataFrame({'food': ['bacon', 'pulled pork', 'bacon', 'Pastrami', 'corned  
      'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})
```



```
In [28]: data
```

```
Out[28]:
```

	food	ounces
0	bacon	4.0
1	pulled pork	3.0
2	bacon	12.0
3	Pastrami	6.0
4	corned beef	7.5
5	Bacon	8.0
6	pastrami	3.0
7	honey ham	5.0
8	nova lox	6.0

```
In [29]: meat_to_animal = {  
    'bacon': 'pig',  
    'pulled pork': 'pig',  
    'pastrami': 'cow',  
    'corned beef': 'cow',  
    'honey ham': 'pig',  
    'nova lox': 'salmon'  
}
```

Here we have a small problem in that some of the meats are capitalized and others are not. Thus, we need to convert each value to lowercase using the `str.lower` Series method:

```
In [31]: lowercased = data['food'].str.lower()
```

```
In [32]: lowercased
```

```
Out[32]: 0      bacon  
1  pulled pork  
2      bacon  
3    pastrami  
4  corned beef  
5      bacon  
6    pastrami  
7    honey ham  
8     nova lox  
Name: food, dtype: object
```

```
In [33]: data['animal'] = lowercased.map(meat_to_animal)
```

```
In [34]: data
```

```
Out[34]:
```

	food	ounces	animal
0	bacon	4.0	pig
1	pulled pork	3.0	pig
2	bacon	12.0	pig
3	Pastrami	6.0	cow
4	corned beef	7.5	cow
5	Bacon	8.0	pig
6	pastrami	3.0	cow
7	honey ham	5.0	pig
8	nova lox	6.0	salmon

## Replacing Values

Filling in missing data with the `fillna` method is a special case of more general value replacement. As you've already seen, `map` can be used to modify a subset of values in an object but `replace` provides a simpler and more flexible way to do so. Let's consider this Series:

```
In [38]: data = pd.Series([1., -999., 2., -999., -1000., 3.])
```

```
In [39]: data
```

```
Out[39]:
```

0	1.0
1	-999.0
2	2.0
3	-999.0
4	-1000.0
5	3.0

dtype: float64

- If you want to replace multiple values at once, you instead pass a list and then the substitute value:

- To use a different replacement for each value, pass a list of substitutes:
- The argument passed can also be a dict:

```
In [40]: data.replace(-999, np.nan)
```

```
Out[40]: 0      1.0  
         1      NaN  
         2      2.0  
         3      NaN  
         4    -1000.0  
         5       3.0  
         dtype: float64
```

```
In [41]: data.replace([-999, -1000], np.nan)
```

```
Out[41]: 0      1.0  
         1      NaN  
         2      2.0  
         3      NaN  
         4      NaN  
         5       3.0  
         dtype: float64
```

```
In [43]: data.replace([-999, -1000], [np.nan, 0])
```

```
Out[43]: 0      1.0  
         1      NaN  
         2      2.0  
         3      NaN  
         4       0.0  
         5       3.0  
         dtype: float64
```

```
In [42]: data.replace({-999: np.nan, -1000: 0})
```

```
Out[42]: 0      1.0  
         1      NaN  
         2      2.0  
         3      NaN  
         4       0.0  
         5       3.0  
         dtype: float64
```

## Renaming Axis Indexes

```
In [46]: data = pd.DataFrame(np.arange(12).reshape((3, 4)),
                             index=['Ohio', 'Colorado', 'New York'],
                             columns=['one', 'two', 'three', 'four'])
```

```
In [47]: data
```

```
Out[47]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
New York	8	9	10	11

**Like a Series, the axis indexes have a map method**

```
In [49]: transform = lambda x: x[:4].upper()
data.index.map(transform)
```

```
Out[49]: Index(['OHIO', 'COLO', 'NEW'], dtype='object')
```

```
In [50]: data.index = data.index.map(transform)
```

```
In [51]: data
```

```
Out[51]:
```

	one	two	three	four
OHIO	0	1	2	3
COLO	4	5	6	7
NEW	8	9	10	11

```
In [55]: data.rename(index=str.title, columns=str.upper)
```

```
Out[55]:
```

	ONE	TWO	THREE	FOUR
Ohio	0	1	2	3
Colo	4	5	6	7
New	8	9	10	11

**Notably, rename can be used in conjunction with a dict-like object providing new values for a subset of the axis labels:**

```
In [ ]: data.rename(index={'OHIO': 'INDIANA'}, columns={'three': 'peekaboo'})
```

**rename saves you from the chore of copying the DataFrame manually and assigning to its**

index and columns attributes. Should you wish to modify a dataset in-place, pass `inplace=True`:

```
In [59]: data.rename(index={'OHIO': 'INDIANA'}, inplace=True)
```

```
In [60]: data
```

```
Out[60]:
```

	one	two	three	four
INDIANA	0	1	2	3
COLO	4	5	6	7
NEW	8	9	10	11

## Discretization and Binning

```
In [63]: ages = [20, 22, 25, 27, 21, 23, 37, 31, 61, 45, 41, 32]
```

Let's divide these into bins of 18 to 25, 26 to 35, 36 to 60, and finally 61 and older. To do so, you have to use `cut`, a function in pandas:

```
In [65]: bins = [18, 25, 35, 60, 100]
cats = pd.cut(ages, bins)
```

```
In [66]: cats
```

```
Out[66]: [(18, 25], (18, 25], (18, 25], (25, 35], (18, 25], ..., (25, 35], (60, 100], (3
5, 60], (35, 60], (25, 35]]
Length: 12
Categories (4, interval[int64]): [(18, 25] < (25, 35] < (35, 60] < (60, 100]]
```

The object pandas returns is a special Categorical object. The output you see describes the bins computed by `pandas.cut`. You can treat it like an array of strings indicating the bin name; internally it contains a categories array specifying the distinct category names along with a labeling for the ages data in the `codes` attribute:

```
In [67]: cats.codes
```

```
Out[67]: array([0, 0, 0, 1, 0, 0, 2, 1, 3, 2, 2, 1], dtype=int8)
```

```
In [68]: cats.categories
```

```
Out[68]: IntervalIndex([(18, 25], (25, 35], (35, 60], (60, 100]],  
                        closed='right',  
                        dtype='interval[int64]')
```

```
In [69]: pd.value_counts(cats)
```

```
Out[69]: (18, 25]      5  
(25, 35]      3  
(35, 60]      3  
(60, 100]     1  
dtype: int64
```

## Detecting and Filtering Outliers

Filtering or transforming outliers is largely a matter of applying array operations. Consider a **DataFrame** with some normally distributed data:

```
In [73]: data = pd.DataFrame(np.random.randn(1000, 4))
```

```
In [74]: data.describe()
```

```
Out[74]:
```

	0	1	2	3
<b>count</b>	1000.000000	1000.000000	1000.000000	1000.000000
<b>mean</b>	0.055569	0.005852	0.006234	-0.057624
<b>std</b>	0.975254	1.048993	0.955249	0.976422
<b>min</b>	-3.383996	-3.452649	-2.916598	-3.118643
<b>25%</b>	-0.581172	-0.690927	-0.655540	-0.742554
<b>50%</b>	0.053096	0.005580	0.033484	-0.084417
<b>75%</b>	0.716883	0.727454	0.636604	0.646606
<b>max</b>	3.145359	3.140651	3.025406	2.836423

Suppose you wanted to find values in one of the columns exceeding 3 in absolute value:

```
In [79]: col = data[2]
```

```
In [80]: col[np.abs(col) > 3]
```

```
Out[80]: 127    3.025406  
         Name: 2, dtype: float64
```

**To select all rows having a value exceeding 3 or –3, you can use the any method on a boolean DataFrame:**

```
In [82]: data[(np.abs(data) > 3).any(1)]
```

```
Out[82]:
```

	0	1	2	3
9	1.182853	3.140651	-0.010715	0.103472
127	-0.573109	-3.200945	3.025406	1.976730
596	-0.024790	-1.537310	1.391747	-3.118643
794	3.145359	-0.843419	1.006886	1.213411
883	0.035233	-3.452649	0.039828	1.309069
987	-3.383996	-0.824807	0.893779	0.988303
991	-0.464648	3.071497	1.982264	-1.456189

**Values can be set based on these criteria. Here is code to cap values outside the interval –3 to 3:**

```
In [83]: data[np.abs(data) > 3] = np.sign(data) * 3
```

```
In [84]: data.describe()
```

```
Out[84]:
```

	0	1	2	3
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	0.055807	0.006293	0.006208	-0.057505
std	0.973523	1.046380	0.955169	0.976057
min	-3.000000	-3.000000	-2.916598	-3.000000
25%	-0.581172	-0.690927	-0.655540	-0.742554
50%	0.053096	0.005580	0.033484	-0.084417
75%	0.716883	0.727454	0.636604	0.646606
max	3.000000	3.000000	3.000000	2.836423

**The statement `np.sign(data)` produces 1 and –1 values based on whether the values in data are positive or negative:**

```
In [88]: data.head()
```

```
Out[88]:
```

	0	1	2	3
0	-0.758658	-0.755808	-0.424850	-1.266991
1	0.860276	-1.892115	-0.954768	0.044466
2	0.195563	-1.523395	0.427470	0.645392
3	-0.530915	-0.079130	0.402178	-0.777428
4	-0.556665	-0.343721	-0.200321	-0.818012

```
In [89]: np.sign(data).head()
```

```
Out[89]:
```

	0	1	2	3
0	-1.0	-1.0	-1.0	-1.0
1	1.0	-1.0	-1.0	1.0
2	1.0	-1.0	1.0	1.0
3	-1.0	-1.0	1.0	-1.0
4	-1.0	-1.0	-1.0	-1.0

## Permutation and Random Sampling

Permuting (randomly reordering) a Series or the rows in a DataFrame is easy to do using the `numpy.random.permutation` function. Calling permutation with the length of the axis you want to permute produces an array of integers indicating the new ordering:

```
In [99]: df = pd.DataFrame(np.arange(5 * 4).reshape((5, 4)))  
sampler = np.random.permutation(5)
```

```
In [100]: df
```

```
Out[100]:
```

	0	1	2	3
0	0	1	2	3
1	4	5	6	7
2	8	9	10	11
3	12	13	14	15
4	16	17	18	19



```
In [101]: sampler
```

```
Out[101]: array([4, 0, 3, 2, 1])
```

```
In [102]: df.take(sampler)
```

```
Out[102]:
```

	0	1	2	3
4	16	17	18	19
0	0	1	2	3
3	12	13	14	15
2	8	9	10	11
1	4	5	6	7

**To select a random subset without replacement, you can use the sample method on Series and DataFrame:**

```
In [105]: df.sample(n=3)
```

```
Out[105]:
```

	0	1	2	3
2	8	9	10	11
0	0	1	2	3
3	12	13	14	15

**To generate a sample with replacement (to allow repeat choices), pass `replace=True` to `sample`:**

```
In [112]: choices = pd.Series([5, 7, -1, 6, 4])  
draws = choices.sample(n=10, replace=True)
```

```
In [113]: draws
```

```
Out[113]:
```

1	7
1	7
3	6
4	4
2	-1
0	5
4	4
3	6
0	5
2	-1

dtype: int64

## Computing Indicator/Dummy Variables

Another type of transformation for statistical modeling or machine learning applications is converting a categorical variable into a “dummy” or “indicator” matrix. If a column in a DataFrame has  $k$  distinct values, you would derive a matrix or DataFrame with  $k$  columns containing all 1s and 0s. pandas has a `get_dummies` function for doing this, though devising one yourself is not difficult. Let’s return to an earlier example DataFrame:

```
In [114]: df = pd.DataFrame({'key': ['b', 'b', 'a', 'c', 'a', 'b'], 'data1': range(6)})
          pd.get_dummies(df['key'])
```

```
Out[114]:
```

	a	b	c
0	0	1	0
1	0	1	0
2	1	0	0
3	0	0	1
4	1	0	0
5	0	1	0

In some cases, you may want to add a prefix to the columns in the indicator DataFrame, which can then be merged with the other data. `get_dummies` has a `prefix` argument for doing this:

```
In [115]: dummies = pd.get_dummies(df['key'], prefix='key')
          df_with_dummy = df[['data1']].join(dummies)
          df_with_dummy
```

```
Out[115]:
```

	data1	key_a	key_b	key_c
0	0	0	1	0
1	1	0	1	0
2	2	1	0	0
3	3	0	0	1
4	4	1	0	0
5	5	0	1	0

A useful recipe for statistical applications is to combine `get_dummies` with a discretization function like `cut`:

```
In [120]: np.random.seed(12345)
values = np.random.rand(10)
```

```
In [121]: values
```

```
Out[121]: array([0.92961609, 0.31637555, 0.18391881, 0.20456028, 0.56772503,
0.5955447 , 0.96451452, 0.6531771 , 0.74890664, 0.65356987])
```

```
In [122]: bins = [0, 0.2, 0.4, 0.6, 0.8, 1]
```

```
In [123]: pd.get_dummies(pd.cut(values, bins))
```

```
Out[123]:
```

	(0.0, 0.2]	(0.2, 0.4]	(0.4, 0.6]	(0.6, 0.8]	(0.8, 1.0]
0	0	0	0	0	1
1	0	1	0	0	0
2	1	0	0	0	0
3	0	1	0	0	0
4	0	0	1	0	0
5	0	0	1	0	0
6	0	0	0	0	1
7	0	0	0	1	0
8	0	0	0	1	0
9	0	0	0	1	0

## String Manipulation

Python has long been a popular raw data manipulation language in part due to its ease of use for string and text processing. Most text operations are made simple with the string object's built-in methods. For more complex pattern matching and text manipulations, regular expressions may be needed. pandas adds to the mix by enabling you to apply string and regular expressions concisely on whole arrays of data, additionally handling the annoyance of missing data.

a comma-separated string can be broken into pieces with split:

```
In [3]: val = 'Mo,Ramdone, Ahmd'
val.split(',')
```

```
Out[3]: ['Mo', 'Ramdone', ' Ahmd']
```

```
In [4]: pieces = [x.strip() for x in val.split(',')]
```

```
In [5]: pieces
```

```
Out[5]: ['Mo', 'Ramdone', 'Ahmd']
```

```
In [7]: name, middle, last = pieces
```

```
In [10]: name + ' ' + middle + ' ' + last
```

```
Out[10]: 'Mo Ramdone Ahmd'
```

```
In [11]: " ".join(pieces)
```

```
Out[11]: 'Mo Ramdone Ahmd'
```

**The difference between index and find, We use both to detect a substring. They both return 1 (True) if found but in case of not finding index raises an exception and find returns (-1).**

```
In [16]: val.index('o')
```

```
Out[16]: 1
```

```
In [18]: val.find('o')
```

```
Out[18]: 1
```

```
In [17]: val.index('z')
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-17-8b62d7293a10> in <module>
----> 1 val.index('z')

ValueError: substring not found
```

```
In [19]: val.find('z')
```

```
Out[19]: -1
```

```
In [20]: val.count(',')
```

```
Out[20]: 2
```

```
In [25]: val.count('m')
```

```
Out[25]: 2
```

```
In [26]: val.replace(',', ' ')
```

```
Out[26]: 'Mo Ramdone  Ahmd'
```

## Regular Expressions

Regular expressions provide a flexible way to search or match (often more complex) string patterns in text. A single expression, commonly called a regex, is a string formed according to the regular expression language. Python's built-in re module is responsible for applying regular expressions to strings; I'll give a number of examples of its use here

```
In [26]: import re
```

The re module functions fall into three categories: pattern matching, substitution, and splitting. Naturally these are all related; a regex describes a pattern to locate in the text, which can then be used for many purposes.

```
In [27]: text = "foo bar\t baz \tqux"
```

we wanted to split a string with a variable number of whitespace characters (tabs, spaces, and newlines). The regex describing one or more whitespace characters is `\s+`:

```
In [28]: re.split('\s+', text)
```

```
Out[28]: ['foo', 'bar', 'baz', 'qux']
```

let's consider a block of text and a regular expression capable of identifying most email addresses:

```
In [29]: text = """Dave dave@google.com
Steve steve@gmail.com
Rob rob@gmail.com
Ryan ryan@yahoo.com
"""
```

```
In [30]: pattern = r'[A-Z0-9._%+-]+@[A-Z0-9.-]+\.[A-Z]{2,4}'
```

```
In [31]: regex = re.compile(pattern, flags=re.IGNORECASE)
```

**Using findall on the text produces a list of the email addresses:**

```
In [32]: regex.findall(text)
```

```
Out[32]: ['dave@google.com', 'steve@gmail.com', 'rob@gmail.com', 'ryan@yahoo.com']
```

```
In [33]: m = regex.search(text)
```

```
In [34]: m
```

```
Out[34]: <re.Match object; span=(5, 20), match='dave@google.com'>
```

```
In [35]: text[m.start():m.end()]
```

```
Out[35]: 'dave@google.com'
```

```
In [36]: pattern = r'([A-Z0-9._%+-]+)@([A-Z0-9.-]+\.[A-Z]{2,4})'
```

```
In [37]: print(regex.sub('REDACTED', text))
```

```
Dave REDACTED
Steve REDACTED
Rob REDACTED
Ryan REDACTED
```

```
In [38]: regex = re.compile(pattern, flags=re.IGNORECASE)
```

```
In [39]: regex.findall(text)
```

```
Out[39]: [('dave', 'google', 'com'),  
          ('steve', 'gmail', 'com'),  
          ('rob', 'gmail', 'com'),  
          ('ryan', 'yahoo', 'com')]
```

```
In [40]: m = regex.match('wesm@bright.net')
```

```
In [41]: m.groups()
```

```
Out[41]: ('wesm', 'bright', 'net')
```

```
In [42]: regex.findall(text)
```

```
Out[42]: [('dave', 'google', 'com'),  
          ('steve', 'gmail', 'com'),  
          ('rob', 'gmail', 'com'),  
          ('ryan', 'yahoo', 'com')]
```

```
In [43]: print(regex.sub(r'Username: \1, Domain: \2, Suffix: \3', text))
```

```
Dave Username: dave, Domain: google, Suffix: com  
Steve Username: steve, Domain: gmail, Suffix: com  
Rob Username: rob, Domain: gmail, Suffix: com  
Ryan Username: ryan, Domain: yahoo, Suffix: com
```

## Vectorized String Functions in pandas

**Cleaning up a messy dataset for analysis often requires a lot of string munging and regularization. To complicate matters, a column containing strings will sometimes have missing data:**

```
In [44]: data = {'Dave': 'dave@google.com', 'Steve': 'steve@gmail.com',  
                'Rob': 'rob@gmail.com', 'Wes': np.nan}
```

```
In [45]: data = pd.Series(data)
```

```
In [46]: data
```

```
Out[46]: Dave      dave@google.com  
Steve    steve@gmail.com  
Rob      rob@gmail.com  
Wes      NaN  
dtype: object
```

```
In [47]: data.isnull()
```

```
Out[47]: Dave      False  
Steve    False  
Rob      False  
Wes      True  
dtype: bool
```

```
In [48]: data.str.contains('gmail')
```

```
Out[48]: Dave      False  
Steve    True  
Rob      True  
Wes      NaN  
dtype: object
```

**You can apply string and regular expression methods can be applied (passing lambda or other function) to each value using `data.map`, but it will fail on the NA (null) values. To cope with this, Series has array-oriented methods for string operations that skip NA values. These are accessed through Series's `str` attribute; for example, we could check whether each email address has 'gmail' in it with `str.contains`:**

**Regular expressions can be used, too, along with any re options like `IGNORECASE`:**

```
In [49]: data.str.findall(pattern, flags=re.IGNORECASE)
```

```
Out[49]: Dave      [(dave, google, com)]  
Steve    [(steve, gmail, com)]  
Rob      [(rob, gmail, com)]  
Wes      NaN  
dtype: object
```

```
In [50]: matches = data.str.match(pattern, flags=re.IGNORECASE)
```



```
In [51]: matches
```

```
Out[51]: Dave      True  
Steve    True  
Rob      True  
Wes      NaN  
dtype: object
```

**We can similarly slice strings using this syntax:**

```
In [54]: data.str[:5]
```

```
Out[54]: Dave      dave@  
Steve    steve  
Rob      rob@g  
Wes      NaN  
dtype: object
```

## Conclusion

Effective data preparation can significantly improve productivity by enabling you to spend more time analyzing data and less time getting it ready for analysis. We have explored a number of tools in this chapter, but the coverage here is by no means comprehensive. In the next chapter, we will explore pandas's joining and grouping functionality.